




(Artificial) Neural Networks: From Perceptron to MLP


Prof. Seungchul Lee
Industrial AI Lab.

Perceptron: Example




$$\begin{aligned}\hat{y} &= g(\omega_0 + X^T \omega) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ &= g(1 + 3x_1 - 2x_2)\end{aligned}$$


Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




Perceptron: Example




$$\hat{y} = g(1 + 3 \times (-1) - 2 \times 2) = g(-6) = -1$$


$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



Perceptron: Forward Propagation




$$\hat{y} = g(\omega_0 + X^T \omega)$$

$$= g\left(\omega_0 + \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}^T \begin{bmatrix} \omega_1 \\ \vdots \\ \omega_m \end{bmatrix}\right)$$

From Perceptron to MLP

Artificial Neural Networks: Perceptron


- Perceptron for $h(\theta)$ or $h(\omega)$
 - Neurons compute the weighted sum of their inputs
 - A neuron is activated or fired when the sum a is positive



- A step function is not differentiable
- One neuron is often not enough
 - One hyperplane

$$a = \omega_0 + \omega_1 x_1 + \omega_2 x_2$$


$$\hat{y} = g(a) = \begin{cases} 1 & a > 0 \\ 0 & \text{otherwise} \end{cases}$$



Here, a step function is illustrated instead of a sign function

XOR Problem


- The main weakness of linear predictors is their lack of capacity.
- For classification, the populations have to be linearly separable.



“xor”




Artificial Neural Networks: MLP

- Multi-layer Perceptron (MLP) = Artificial Neural Networks (ANN)
 - Multi neurons = multiple linear classification boundaries




Artificial Neural Networks: Activation Function

- Differentiable nonlinear activation function




Artificial Neural Networks

- In a compact representation



Artificial Neural Networks

- A single layer is not enough to be able to represent complex relationship between input and output
⇒ perceptron with many layers and units




- Multi-layer perceptron
 - Features of features
 - Mapping of mappings

Another Perspective: ANN as Kernel Learning

Nonlinear Mapping


- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.


$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.


$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Neuron

- We can represent this “neuron” as follows:


$$f(x) = \sigma(w \cdot x + b)$$



Kernel + Neuron


- Nonlinear mapping + neuron

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Neuron + Neuron


- Nonlinear mapping can be represented by another neurons



- Nonlinear Kernel
 - Nonlinear activation functions



Multi Layer Perceptron

- Nonlinear mapping can be represented by another neurons
- We can generalize an MLP




Summary

- Universal function approximator
- Universal function classifier
- Parameterized




Artificial Neural Networks

- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons




Deep Artificial Neural Networks

- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons





Looking at Parameters

Logistic Regression in a Form of Neural Network




$$y = \sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2)$$




Logistic Regression in a Form of Neural Network

- Neural network convention


$$y = \sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2)$$




$$y = \sigma(b + \omega_1 x_1 + \omega_2 x_2)$$



Do not indicate bias units





```
n_input = 2  
n_output = 1
```



Logistic Regression in a Form of Neural Network


- One-hot encoding
 - One-hot encoding is a conventional practice for a multi-class classification

$$y^{(i)} \in \{1, 0\} \implies y^{(i)} \in \{[0, 1], [1, 0]\}$$




Nonlinearly Distributed Data


- Example to understand network's behavior
 - Include a hidden layer



Do not include bias units




```
n_input = 2  
n_hidden = 2  
n_output = 2
```




Multi Layers


- x space



Do not include bias units




```
n_input = 2  
n_hidden = 2  
n_output = 2
```




Multi Layers


- z space



Do not include bias units




```
n_input = 2  
n_hidden = 2  
n_output = 2
```




Nonlinearly Distributed Data


- More neurons in hidden layer



Do not include bias units




```
n_input = 2  
n_hidden = 4  
n_output = 2
```




Multi Layers


- Multiple linear classification boundaries



Do not include bias units



```
n_input = 2  
n_hidden = 4  
n_output = 2
```



(Artificial) Neural Networks: Training

Training Neural Networks: Loss Function

- Measures error between target values and predictions

$$\min_{\omega} \sum_{i=1}^m \ell \left(h_{\omega} \left(x^{(i)} \right), y^{(i)} \right)$$

- Example

- Squared loss (for regression):

$$\frac{1}{m} \sum_{i=1}^m \left(h_{\omega} \left(x^{(i)} \right) - y^{(i)} \right)^2$$

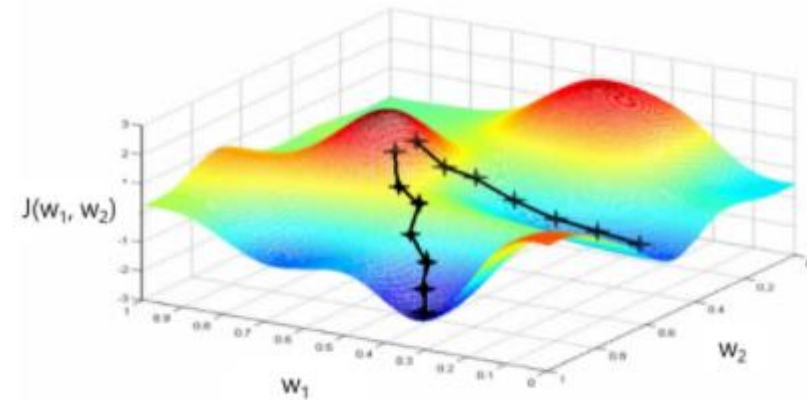
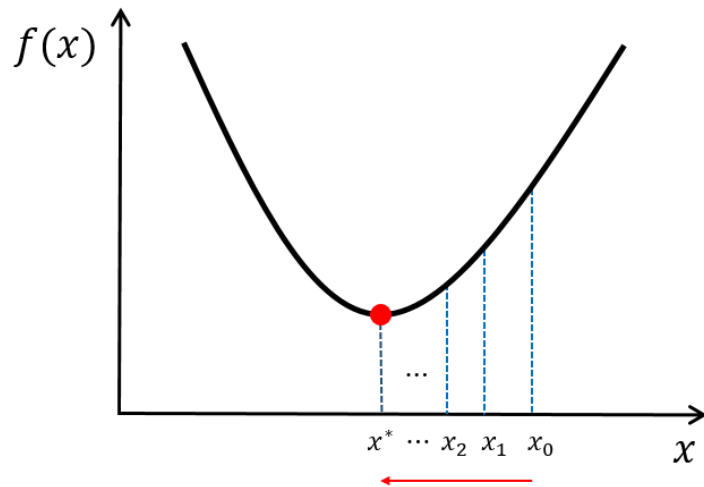
- Cross entropy (for classification):

$$-\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \left(h_{\omega} \left(x^{(i)} \right) \right) + \left(1 - y^{(i)} \right) \log \left(1 - h_{\omega} \left(x^{(i)} \right) \right)$$

Training Neural Networks: Gradient Descent



- Negative gradients points directly downhill of the cost function
- We can decrease the cost by moving in the direction of the negative gradient (α is a learning rate)

$$\omega \leftarrow \omega - \alpha \nabla_{\omega} \ell \left(h_{\omega} \left(x^{(i)} \right), y^{(i)} \right)$$



Gradients in ANN


- Learning weights and biases from data using gradient descent
- $\frac{\partial \ell}{\partial \omega}$: too many computations are required for all ω
- Structural constraint of NN:
 - Composition of functions
 - Chain rule
 - Dynamic programming




ANN in TensorFlow: MNIST

MNIST database

- Mixed National Institute of Standards and Technology database
- Handwritten digit database
- 28×28 gray scaled image
- **Flattened** matrix into a vector of $28 \times 28 = 784$



Our Network Model



Implementation in Python

```
mnist = tf.keras.datasets.mnist

(train_x, train_y), (test_x, test_y) = mnist.load_data()


train_x, test_x = train_x/255.0, test_x/255.0
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape = (28, 28)),
    tf.keras.layers.Dense(units = 100, activation = 'relu'),
    tf.keras.layers.Dense(units = 10, activation = 'softmax')
])

model.compile(optimizer = 'adam',
              loss = 'sparse_categorical_crossentropy',
              metrics = ['accuracy'])

loss = model.fit(train_x, train_y, epochs = 5)
```

```
test_loss, test_acc = model.evaluate(test_x, test_y)
```



Evaluation

```
test_img = test_x[np.random.choice(test_x.shape[0], 1)]  
  
predict = model.predict_on_batch(test_img)  
mypred = np.argmax(predict, axis = 1)
```

